


CSA vs TWFE


273 retained / 369 comparisons; 32 capped (diamonds)



- Study 1 (2017)
- Study 2 (2017)
- Study 4 (2018)
- Study 5 (2018)
- Study 6 (2018)
- Study 7 (2019)
- Study 8 (2020)
- Study 9 (2020)
- Study 10 (2021)
- Study 11 (2021)
- Study 12 (2021)
- Study 13 (2022)

BJS vs TWFE

297 retained / 373 comparisons; 43 capped (diamonds)



- Study 14 (2022)
- Study 15 (2023)
- Study 16 (2023)
- Study 17 (2023)
- Study 18 (2024)
- Study 19 (2024)
- Study 20 (2024)
- Study 22 (2024)
- Study 23 (2025)
- Study 24 (2025)

Notes: Circles: static specifications. Squares: event-study averages, with pre and post entering separately. Pre averages are diagnostics, not treatment effects. Shares use retained comparisons, including flagged points. Significant + Reversal: opposite nonzero signs and a filled marker. Diamonds: coordinates capped at ± 5 ; displayed gaps may differ from actual gaps. Excluded if either estimate has effective clusters < 5 or maximum cluster leverage $> 50\%$. Crosses: effective clusters < 10 or maximum leverage $> 25\%$ for either estimate or their difference, or prior inference warnings; significance fill suppressed. Leverage uses squared outcome-linear weights by the existing bootstrap clusters; it is a design diagnostic, not a variance guarantee. TWFE retains its eligible sample; modern estimators use their own supported targets. Both axes divide by this pair's saved TWFE SE; the plotted gap is not the paired-difference t statistic. Common horizons are unchanged. Fill uses paired-bootstrap normal tests, without multiplicity adjustment. Dashed line: equality. Study 8 uses unrestricted TWFE; its restricted published fit is retained separately. Study 4 clusters 143 builders despite only three treatment-rollout regions.